

Reinforcement Learning

Shusen Wang

A little bit probability theory...

Random Variable

- **Random variable**: unknown; its values depends on outcomes of a random event.
- Uppercase letter ***X*** for random variable.

*Random
Variable*

*Possible
Values*

*Random
Events*

Probabilities

$$X = \begin{cases} 0 \\ 1 \end{cases}$$



$$\mathbb{P}(X = 0) = 0.5$$

$$\mathbb{P}(X = 1) = 0.5$$

Random Variable

- **Random variable**: unknown; its values depends on outcomes of a random event.
- Uppercase letter X for random variable.
- Lowercase letter x for an observed value.
- For example, I flipped a coin 4 times and observed:
 - $x_1 = 1$,
 - $x_2 = 1$,
 - $x_3 = 0$,
 - $x_4 = 1$.

Probability Density Function (PDF)

- PDF provides a relative likelihood that the value of the random variable would equal that sample.


Probability Density Function (PDF)

- PDF provides a relative likelihood that the value of the random variable would equal that sample.

Example: Gaussian distribution

- It is a continuous distribution.
- PDF:

$$p(x) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(x-\mu)^2}{2\sigma^2}\right).$$



Probability Density Function (PDF)

- PDF provides a relative likelihood that the value of the random variable would equal that sample.


Example: Discrete distribution

- Discrete random variable: $X \in \{1, 3, 7\}$.
- PDF:

$$p(1) = 0.2,$$

$$p(3) = 0.5,$$

$$p(7) = 0.3.$$



Probability Density Function (PDF)

- Random variable X is in the domain \mathcal{X} .
- For continuous distribution,

$$\int_{\mathcal{X}} p(x) dx = 1.$$

- For discrete distribution,

$$\sum_{x \in \mathcal{X}} p(x) = 1.$$

Expectation

- Random variable X is in the domain \mathcal{X} .
- For continuous distribution, the expectation of $f(X)$ is:


$$\mathbb{E} [f(X)] = \int_{\mathcal{X}} p(x) \cdot f(x) dx.$$

- For discrete distribution, the expectation of $f(X)$ is:

$$\mathbb{E} [f(X)] = \sum_{x \in \mathcal{X}} p(x) \cdot f(x) .$$

Random Sampling

- There are 10 balls in the bin: 2 are red, 5 are green, and 3 are blue.
- Randomly sample a ball.
- What will be the color?



Random Sampling

- Sample red ball w.p. 0.2, green ball w.p. 0.5, and blue ball w.p. 0.3.
- Randomly sample a ball.
- What will be the color?

Random Sampling

- Sample **red ball w.p. 0.2**, **green ball w.p. 0.5**, and **blue ball w.p. 0.3**.
- Randomly sample a ball.
- What will be the color?

```
from numpy.random import choice
```

```
samples = choice(['R', 'G', 'B'], size=100, p=[0.2, 0.5, 0.3])
```

```
print(samples)
```


```
[ 'R'  'G'  'R'  'R'  'R'  'R'  'B'  'B'  'B'  'G'  'G'  'B'  'G'  'B'  'B'  'G'  'B'  'G'
  'B'  'B'  'G'  'B'  'G'  'B'  'B'  'G'  'B'  'B'  'G'  'B'  'G'  'G'  'G'  'G'  'G'  'B'
  'B'  'B'  'B'  'B'  'B'  'G'  'G'  'B'  'R'  'R'  'B'  'R'  'B'  'G'  'R'  'G'  'R'  'G'
  'R'  'R'  'B'  'G'  'G'  'G'  'B'  'R'  'G'  'B'  'G'  'R'  'G'  'G'  'G'  'B'  'B'  'R'
  'G'  'G'  'B'  'B'  'R'  'B'  'B'  'B'  'R'  'B'  'G'  'B'  'R'  'B'  'R'  'G'  'B'  'R'
  'B'  'B'  'G'  'G'  'G'  'R'  'R'  'B'  'R'  'G' ]
```

Terminologies

Terminology: state and action

state s (this frame)

Action $a \in \{\text{left}, \text{right}, \text{up}\}$




Terminology: state and action

state s (this frame)




Action $a \in \{\text{left}, \text{right}, \text{up}\}$



Terminology: policy

policy π




- Policy function $\pi: (s, a) \mapsto [0,1]$:
$$\pi(a | s) = \mathbb{P}(A = a | S = s).$$
- It is the probability of taking action $A = a$ given state s , e.g.,
 - $\pi(\text{left} | s) = 0.2,$
 - $\pi(\text{right} | s) = 0.1,$
 - $\pi(\text{up} | s) = 0.7.$
- Upon observing state $S = s$, the agent's action A can be random.

Terminology: policy

policy π


- Policy function $\pi: (s, a) \mapsto [0,1]$:
$$\pi(a | s) = \mathbb{P}(A = a | S = s).$$
- It is the probability of taking action $A = a$ given state s , e.g.,
 - $\pi(\text{left} | s) = 0.2$,
 - $\pi(\text{right} | s) = 0.1$,
 - $\pi(\text{up} | s) = 0.7$.
- Upon observing state $S = s$, the agent's action A can be random.



Terminology: policy


policy π

- Policy function $\pi: (s, a) \mapsto [0,1]$:
$$\pi(a | s) = \mathbb{P}(A = a | S = s).$$
- It is the probability of taking action $A = a$ given state s , e.g.,
 - $\pi(\text{left} | s) = 0.2$,
 - $\pi(\text{right} | s) = 0.1$,
 - $\pi(\text{up} | s) = 0.7$.
- Upon observing state $S = s$, the agent's action A can be random.



Terminology: policy

policy π



- Policy function $\pi: (s, a) \mapsto [0,1]$:
$$\pi(a | s) = \mathbb{P}(A = a | S = s).$$
- It is the probability of taking action $A = a$ given state s , e.g.,
 - $\pi(\text{left} | s) = 0.2$,
 - $\pi(\text{right} | s) = 0.1$,
 - $\pi(\text{up} | s) = 0.7$.
- Upon observing state $S = s$, the agent's action A can be random.

Terminology: reward

reward R

- Collect a coin: $R = +1$



Terminology: reward

reward R



- Collect a coin: $R = +1$
- Win the game: $R = +10000$

Terminology: reward

reward R



- Collect a coin: $R = +1$
- Win the game: $R = +10000$
- Touch a Goomba: $R = -10000$ (game over).

Terminology: reward

reward R



- Collect a coin: $R = +1$
- Win the game: $R = +10000$
- Touch a Goomba: $R = -10000$ (game over).
- Nothing happens: $R = 0$

Terminology: state transition




state transition




Terminology: state transition

state transition



- E.g., “up” action leads to a new state.

Terminology: state transition




state transition



- E.g., “up” action leads to a new state.
- State transition can be random.
- Randomness is from the environment.

Terminology: state transition




state transition

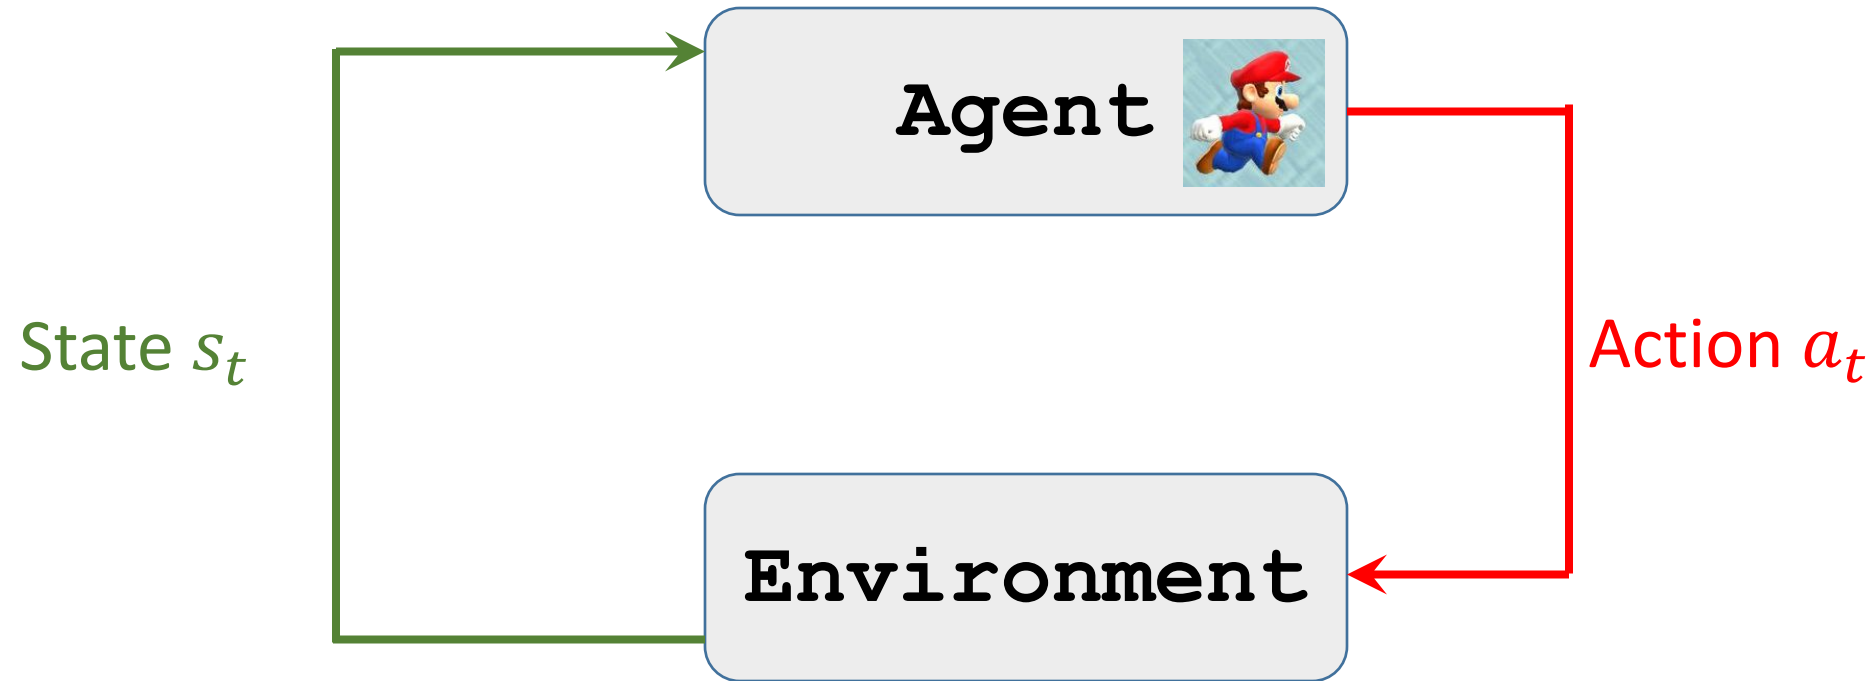


- E.g., “up” action leads to a new state.
- State transition can be random.
- Randomness is from the environment.
- $p(s'|s, a) = \mathbb{P}(S' = s' | S = s, A = a)$.

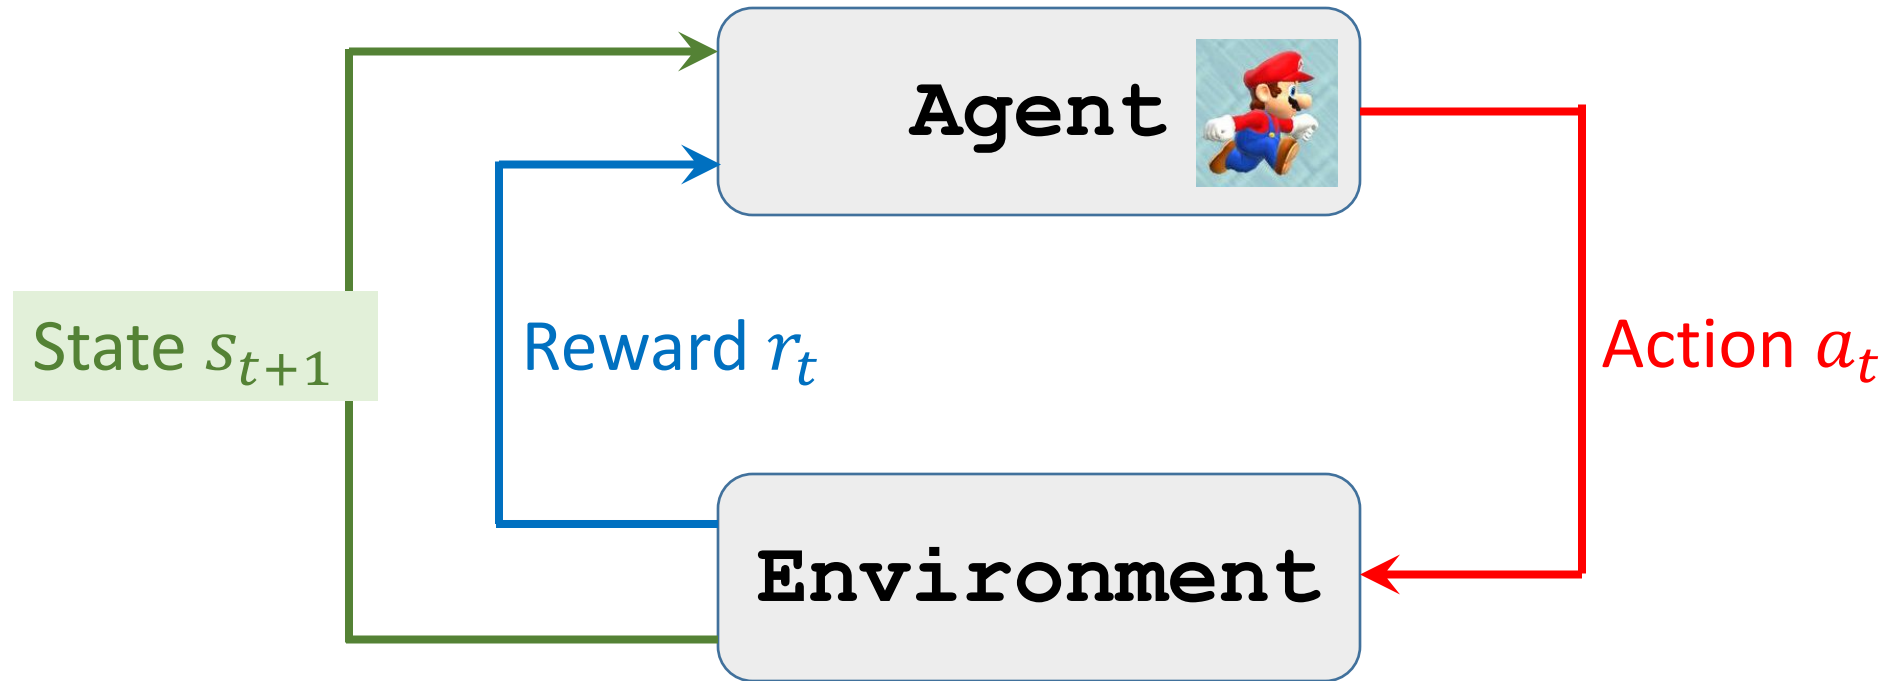
Terminology: agent environment interaction



Terminology: agent environment interaction



Terminology: agent environment interaction




Randomness in Reinforcement Learning


Actions have randomness.

- Given state s , the action can be random, e.g., .

- $\pi(\text{"left"}|s) = 0.2,$
- $\pi(\text{"right"}|s) = 0.1,$
- $\pi(\text{"up"}|s) = 0.7.$



Randomness in Reinforcement Learning




Actions have randomness.

- Given state s , the action can be random, e.g., .
 - $\pi(\text{"left"}|s) = 0.2$,
 - $\pi(\text{"right"}|s) = 0.1$,
 - $\pi(\text{"up"}|s) = 0.7$.

State transitions have randomness.

- Given state $S = s$ and action $A = a$, the environment randomly generates a new state S' .

Randomness in Reinforcement Learning



Actions have randomness.

- Given state s , the action can be random, e.g., .
 - $\pi(\text{"left"}|s) = 0.2$,
 - $\pi(\text{"right"}|s) = 0.1$,
 - $\pi(\text{"up"}|s) = 0.7$.

State transitions have randomness.

- Given state $S = s$ and action $A = a$, the environment randomly generates a new state S' .

Play the game using AI



- Observe a frame (state s_1)
- ➔ Make action a_1 (left, right, or up)
- ➔ Observe a new frame (state s_2) and reward r_1
- ➔ Make action a_2
- ➔ ...

Play the game using AI



- Observe a frame (state s_1)
- ➔ Make action a_1 (left, right, or up)
- ➔ Observe a new frame (state s_2) and reward r_1
- ➔ Make action a_2
- ➔ ...
- (state, action, reward) trajectory:
 $s_1, a_1, r_1, s_2, a_2, r_2, \dots, s_T, a_T, r_T.$

Rewards and Returns

Return

Definition: Return (aka cumulative future reward).

- $U_t = R_t + R_{t+1} + R_{t+2} + R_{t+3} + \dots$

Return

Definition: Return (aka cumulative future reward).

- $U_t = R_t + R_{t+1} + R_{t+2} + R_{t+3} + \dots$

Question: Are R_t and R_{t+1} equally important?

- Which of the followings do you prefer?
 - I give you \$100 right now.
 - I will give you \$100 one year later.

Return

Definition: Return (aka cumulative future reward).

- $U_t = R_t + R_{t+1} + R_{t+2} + R_{t+3} + \dots$

Question: Are R_t and R_{t+1} equally important?

- Which of the followings do you prefer?
 - I give you \$100 right now.
 - I will give you \$100 one year later.
- Future reward is less valuable than present reward.
- R_{t+1} should be given less weight than R_t .

Return

Definition: Return (aka cumulative future reward).

- $U_t = R_t + R_{t+1} + R_{t+2} + R_{t+3} + \dots$

Definition: Discounted return (aka cumulative discounted future reward).

- γ : discount rate (tuning hyper-parameter).
- $U_t = R_t + \gamma R_{t+1} + \gamma^2 R_{t+2} + \gamma^3 R_{t+3} + \dots$

Randomness in Returns

Definition: Discounted return (at time step t).

- $U_t = R_t + \gamma R_{t+1} + \gamma^2 R_{t+2} + \gamma^3 R_{t+3} + \dots$

At time step t , the return U_t is **random**.

- Two sources of randomness:

1. Action can be random: $\mathbb{P}[A = a \mid S = s] = \pi(a|s)$.
2. New state can be random: $\mathbb{P}[S' = s' \mid S = s, A = a] = p(s'|s, a)$.

Randomness in Returns

Definition: Discounted return (at time step t).

- $U_t = R_t + \gamma R_{t+1} + \gamma^2 R_{t+2} + \gamma^3 R_{t+3} + \dots$

At time step t , the return U_t is **random**.

- Two sources of randomness:
 1. Action can be random: $\mathbb{P}[A = a \mid S = s] = \pi(a|s)$.
 2. New state can be random: $\mathbb{P}[S' = s' \mid S = s, A = a] = p(s'|s, a)$.
- For any $i \geq t$, the reward R_i depends on S_i and A_i .
- Thus, given s_t , the return U_t depends on the random variables:
 - $A_t, A_{t+1}, A_{t+2}, \dots$ and S_{t+1}, S_{t+2}, \dots .

Value Functions

Action-Value Function $Q(s, a)$

Definition: Discounted return (aka cumulative discounted future reward).

- $U_t = R_t + \gamma R_{t+1} + \gamma^2 R_{t+2} + \gamma^3 R_{t+3} + \dots$

Action-Value Function $Q(s, a)$

Definition: Discounted return (aka cumulative discounted future reward).

- $U_t = R_t + \gamma R_{t+1} + \gamma^2 R_{t+2} + \gamma^3 R_{t+3} + \dots$

Definition: Action-value function for policy π .

- $Q_\pi(s_t, a_t) = \mathbb{E} [U_t | S_t = s_t, A_t = a_t].$



- Return U_t depends on states $s_t, s_{t+1}, s_{t+2}, \dots$ and actions $a_t, a_{t+1}, a_{t+2}, \dots$.

Action-Value Function $Q(s, a)$

Definition: Discounted return (aka cumulative discounted future reward).

- $U_t = R_t + \gamma R_{t+1} + \gamma^2 R_{t+2} + \gamma^3 R_{t+3} + \dots$

Definition: Action-value function for policy π .

- $Q_\pi(s_t, a_t) = \mathbb{E} [U_t | S_t = s_t, A_t = a_t].$



- Return U_t depends on states $s_t, s_{t+1}, s_{t+2}, \dots$ and actions $a_t, a_{t+1}, a_{t+2}, \dots$.
- Actions are random: $\mathbb{P}[A = a | S = s] = \pi(a|s).$ (Policy function.)
- States are random: $\mathbb{P}[S' = s' | S = s, A = a] = p(s'|s, a).$ (State transition.)

Action-Value Function $Q(s, a)$

Definition: Discounted return (aka cumulative discounted future reward).

- $U_t = R_t + \gamma R_{t+1} + \gamma^2 R_{t+2} + \gamma^3 R_{t+3} + \dots$

Definition: Action-value function for policy π .

- $Q_\pi(s_t, a_t) = \mathbb{E} [U_t | S_t = s_t, A_t = a_t].$

Definition: Optimal action-value function.

- $Q^*(s_t, a_t) = \max_{\pi} Q_\pi(s_t, a_t).$

State-Value Function $V(s)$

Definition: Discounted return (aka cumulative discounted future reward).

- $U_t = R_t + \gamma R_{t+1} + \gamma^2 R_{t+2} + \gamma^3 R_{t+3} + \dots$

Definition: Action-value function for policy π .

- $Q_\pi(s_t, a_t) = \mathbb{E} [U_t | S_t = s_t, A_t = a_t].$

Definition: State-value function.

- $V_\pi(s_t) = \mathbb{E}_A [Q_\pi(s_t, A)]$

State-Value Function $V(s)$

Definition: Discounted return (aka cumulative discounted future reward).

- $U_t = R_t + \gamma R_{t+1} + \gamma^2 R_{t+2} + \gamma^3 R_{t+3} + \dots$

Definition: Action-value function for policy π .

- $Q_\pi(s_t, a_t) = \mathbb{E} [U_t | S_t = s_t, A_t = a_t].$

Definition: State-value function.

- $V_\pi(s_t) = \mathbb{E}_A [Q_\pi(s_t, A)] = \sum_a \pi(a | s_t) \cdot Q_\pi(s_t, a). \quad (\text{Actions are discrete.})$

Taken w.r.t. the action $A \sim \pi(\cdot | s_t).$

State-Value Function $V(s)$

Definition: Discounted return (aka cumulative discounted future reward).

- $U_t = R_t + \gamma R_{t+1} + \gamma^2 R_{t+2} + \gamma^3 R_{t+3} + \dots$

Definition: Action-value function for policy π .

- $Q_\pi(s_t, a_t) = \mathbb{E} [U_t | S_t = s_t, A_t = a_t].$

Definition: State-value function.

- $V_\pi(s_t) = \mathbb{E}_A [Q_\pi(s_t, A)] = \sum_a \pi(a|s_t) \cdot Q_\pi(s_t, a).$ (Actions are discrete.)
- $V_\pi(s_t) = \mathbb{E}_A [Q_\pi(s_t, A)] = \int \pi(a|s_t) \cdot Q_\pi(s_t, a) da.$ (Actions are continuous.)

Understanding the Value Functions

- **Action**-value function: $Q_{\pi}(s_t, a_t) = \mathbb{E} [U_t | S_t = s_t, A_t = a_t]$.
- Given policy π , $Q_{\pi}(s, a)$ evaluates how good it is for an agent to pick action a while being in state s .

Understanding the Value Functions

- **Action**-value function: $Q_{\pi}(s_t, a_t) = \mathbb{E} [U_t | S_t = s_t, A_t = a_t]$.
- Given policy π , $Q_{\pi}(s, a)$ evaluates how good it is for an agent to pick action a while being in state s .
- **State**-value function: $V_{\pi}(s) = \mathbb{E}_A [Q_{\pi}(s, A)]$
- For fixed policy π , $V_{\pi}(s)$ evaluates how good the situation is in state s .
- $\mathbb{E}_S[V_{\pi}(S)]$ evaluates how good the policy π is.

Play games using reinforcement learning

How does AI control the agent?

Suppose we have a good policy $\pi(a|s)$.

- Upon observing the state s_t ,
- random sampling: $a_t \sim \pi(\cdot | s_t)$.

How does AI control the agent?

Suppose we have a good policy $\pi(a|s)$.

- Upon observing the state s_t ,
- random sampling: $a_t \sim \pi(\cdot | s_t)$.


Suppose we know the optimal action-value function $Q^*(s, a)$.

- Upon observe the state s_t ,
- choose the **action** that maximizes the value: $a_t = \operatorname{argmax}_a Q^*(s_t, a)$.

OpenAI Gym

- Gym is a toolkit for developing and comparing reinforcement learning algorithms.
- <https://gym.openai.com/>

Classical control problems



Cart Pole




Pendulum

OpenAI Gym

- Gym is a toolkit for developing and comparing reinforcement learning algorithms.
- <https://gym.openai.com/>

Atari Games



Pong



Space Invader




Breakout


OpenAI Gym

- Gym is a toolkit for developing and comparing reinforcement learning algorithms.
- <https://gym.openai.com/>


MuJoCo (Continuous control tasks.)



Ant




Humanoid



Half Cheetah


OpenAI Gym



Play CartPole Game

```
import gym
env = gym.make( 'CartPole-v0' )
```

- Get the environment of CartPole from Gym.
- “env” provides states and reward.



Play CartPole Game

```
state = env.reset()
```

```
for t in range(100):
```

```
    env.render()
```

```
    print(state)
```

A window pops up rendering CartPole.

A random **action**.

```
    action = env.action_space.sample()
```

```
    state, reward, done, info = env.step(action)
```

```
    if done: "done=1" means finished (win or lose the game)
```

```
        print('Finished')
```

```
        break
```

```
env.close()
```

Summary

Summary

Terminologies

- Agent 
- Environment
- State s .
- Action a .
- Reward r .
- Policy $\pi(a|s)$
- State transition $p(s'|s, a)$.

Summary

Terminologies

- Agent 
- Environment
- State s .
- Action a .
- Reward r .
- Policy $\pi(a|s)$
- State transition $p(s'|s, a)$.

Return and Value

- Return:

$$U_t = R_t + \gamma R_{t+1} + \gamma^2 R_{t+2} + \dots$$

- Action-value function:

$$Q_\pi(s_t, a_t) = \mathbb{E}[U_t | s_t, a_t].$$

- Optimal action-value function:


$$Q^*(s_t, a_t) = \max_{\pi} Q_\pi(s_t, a_t).$$

- State-value function:

$$V_\pi(s_t) = \mathbb{E}_A[Q_\pi(s_t, A)].$$

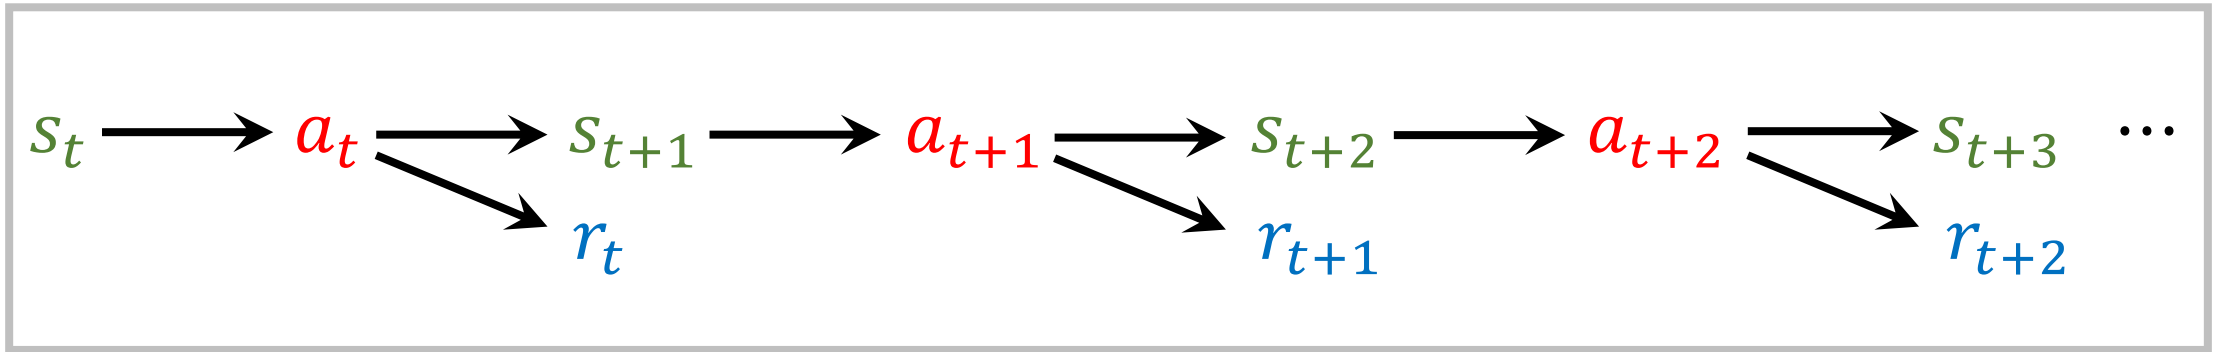
Play game using reinforcement learning

- Observe state s_t , make action a_t , environment gives s_{t+1} and reward r_t .



Play game using reinforcement learning

- Observe state s_t , make action a_t , environment gives s_{t+1} and reward r_t .



- The agent can be controlled by either $\pi(a|s)$ or $Q^*(s, a)$.

We are going to study...

2. Value-based learning.

- Deep Q network (DQN) for approximating $Q^*(s, a)$.
- Learn the network parameters using temporal different (TD).

3. Policy-based learning.

- Policy network for approximating $\pi(a|s)$.
- Learn the network parameters using policy gradient.

4. Actor-critic method. (Policy network + value network.)

5. Example: AlphaGo

Thank you!